

Class 08: Unsupervised Learning Mini-project

Heidi Nam

Exploratory Data Analysis

We are using the excel file "WisconsinCancer.csv" as our data. We will download and read the csv file:

```
fna.data <- "WisconsinCancer.csv"
wisc.df <- read.csv(fna.data, row.names=1)
```

As we don't need the diagnosis table within our data frame `wisc.df`, we will record it separately into a `diagnosis` vector.

```
wisc.data <- wisc.df[, -1]
diagnosis <- factor(wisc.df[, 1])
```

Answering questions 1~3:

Q1: How many observations are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

There are 569 observations in this dataset.

Q2. How many of the observations have a malignant diagnosis?

```
sum(diagnosis == "M")
```

```
[1] 212
```

There are 212 observations with a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10 variables that are suffixed with `"_mean"`.

Principal Component Analysis (PCA)

Checking if the data needs to be scaled before performing PCA:

```
# Check column means and standard deviations  
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Performing PCA with scaling:

```
wisc.pr <- prcomp(wisc.data,scale. = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005

Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Answering questions 4-6:

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

A: 0.4427 (44.27%) of the original variance is captured by the first principal component PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

A: The first three principal components describe at least 70% of the original variance in the data.


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

A: The first seven principal components describe at least 90% of the original variance.

Interpreting PCA results

Plotting the PCA results into a biplot:

```
biplot(wisc.pr)
```

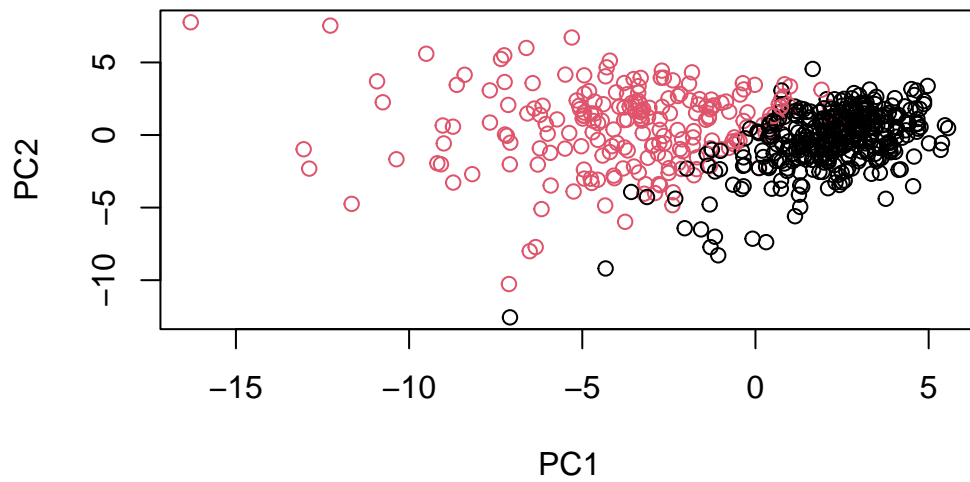


Q7: What stands out to you about this plot? Is it easy or difficult to understand? Why?

A: This plot is able to plot both the diagnosis number and the different variables and how they can be expressed on the plot with the first and second principal components. However, This is very difficult to understand as we don't understand why the plot originates from one center origin (and what it signifies) as well as have a lot of overlap of text and lines that makes the plot difficult to read.

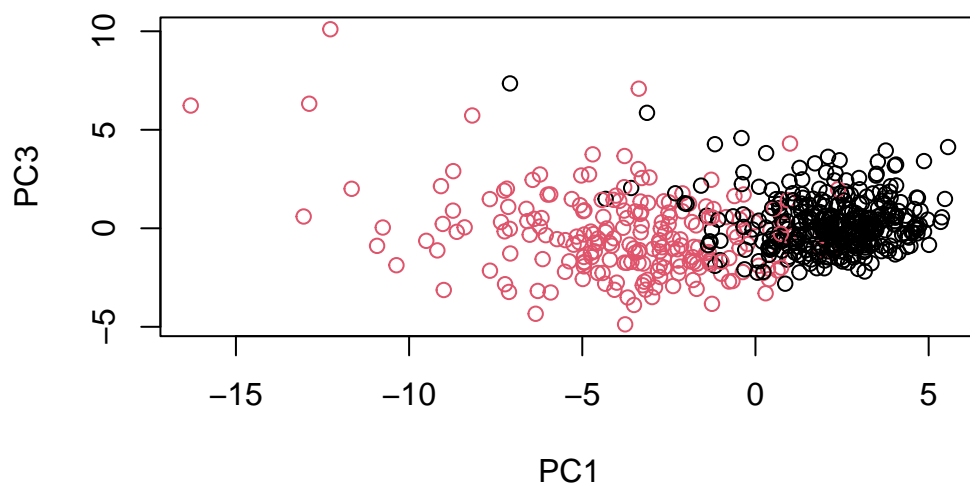
Instead, plotting a scatter plot to observe the relation of the two principal components that are categorized by their diagnosis (malignant or benign).

```
plot( wisc.pr$x[,1:2], col=diagnosis,
      xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot( wisc.pr$x[,c(1,3)], col=diagnosis,  
      xlab = "PC1", ylab = "PC3")
```



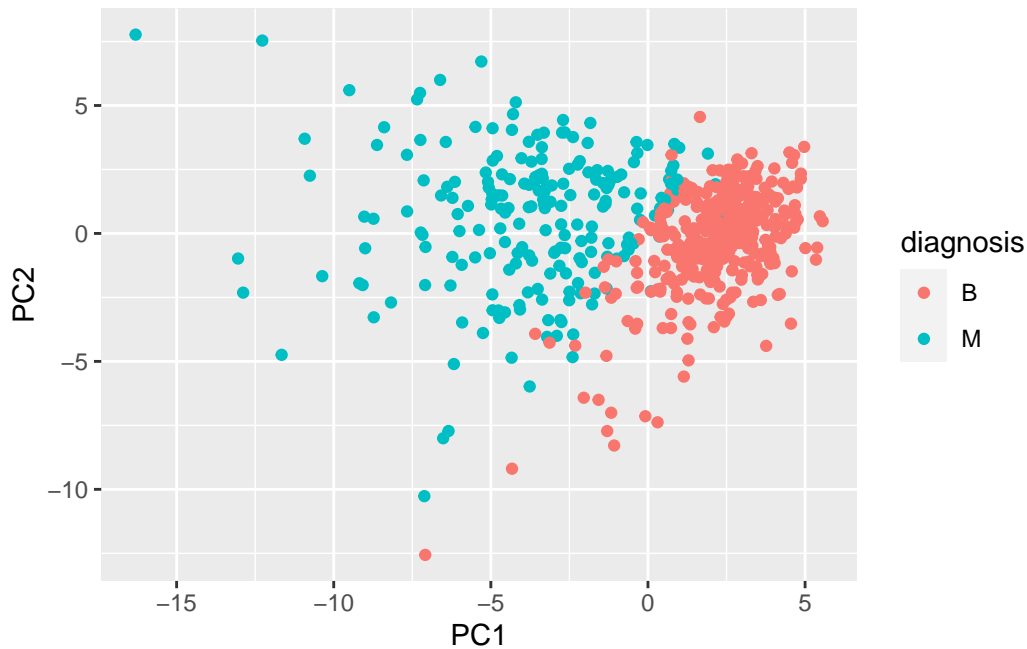
The plot feels a lot more cluttered and its outliers feel a lot more variant as it has a higher range of PC3; thus, it feels like it doesn't optimally capture the data well.

Using ggplot to understand the data in a more visually aesthetic platform:

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

library(ggplot2)

ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Variance explained

Understanding the proportion of variance by calculating the variance of each principal component by taking its standard deviation and squaring it:

```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

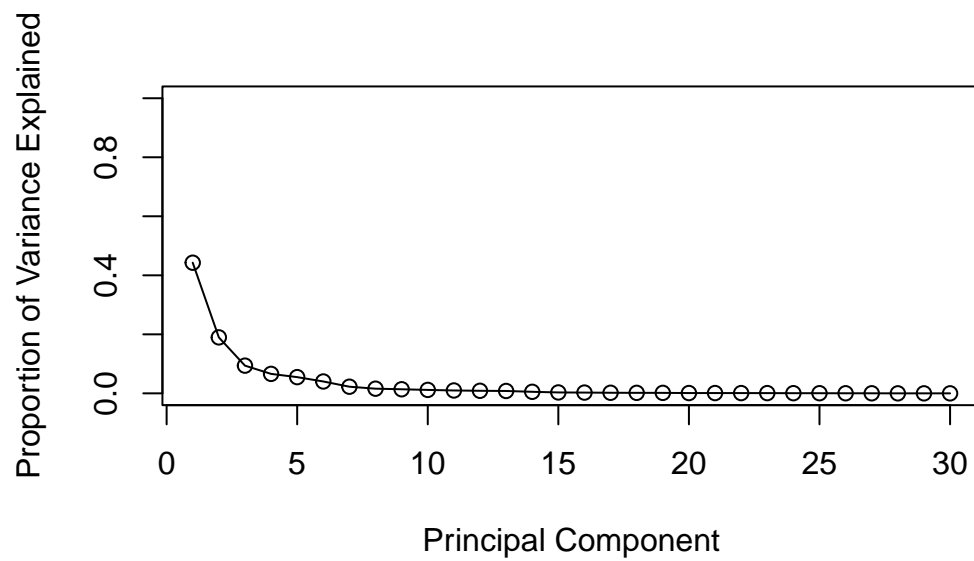
```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Calculating proportion of variance by dividing the variance by total variance of principal components:

```
pve <- pr.var / sum(pr.var)
```


Plotting this into a scree plot:

```
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```

Presenting the same data in a bar plot:

```
barplot(pve, ylab = "Precent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```

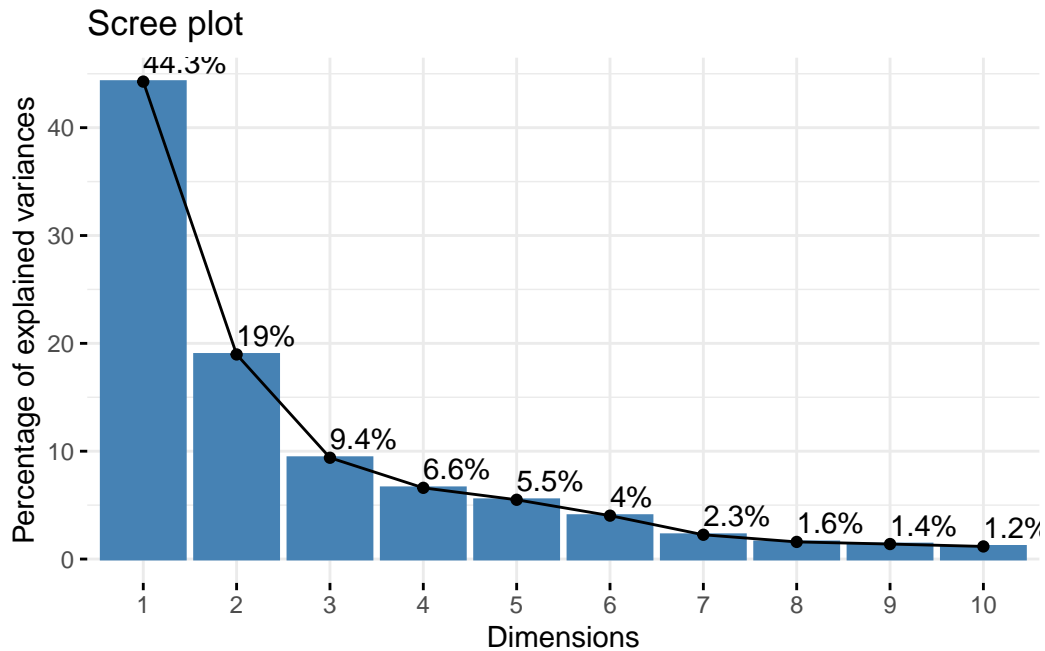


Plotting using a ggplot based graph:

```
# install.packages("factoextra")  
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

```
loading_vector <- wisc.pr$rotation[,1]
loading_vector["concave.points_mean"]
```

```
concave.points_mean
-0.2608538
```

The component of the loading vector is -0.2608538.

Hierarchical Clustering

Scaling `wisc.data` by using the `scale()` function:

```
data.scaled <- scale(wisc.data)
```

Calculating distances between all pairs of observations:


```
data.dist <- dist(data.scaled)
```

Creating hierarchical clustering model using complete linkage:

```
wisc.hclust <- hclust(data.dist, "complete")
```

Q10: Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```



```
data.dist  
hclust (*, "complete")
```

A: The height is 19 in where the clustering model has 4 clusters.

Selecting number of clusters

Cutting the cluster tree to have 4 clusters:

```
wisc.hclust.clusters <- cutree(wisc.hclust,4)
```

Comparing cluster membership to the diagnoses:

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.hclust.clusters	B	M	
1	12	165	
2	2	5	
3	343	40	
4	0	2	

Using different methods

Q12. Which method gives your favorite results for the same `data.dist` dataset? Explain your reasoning.

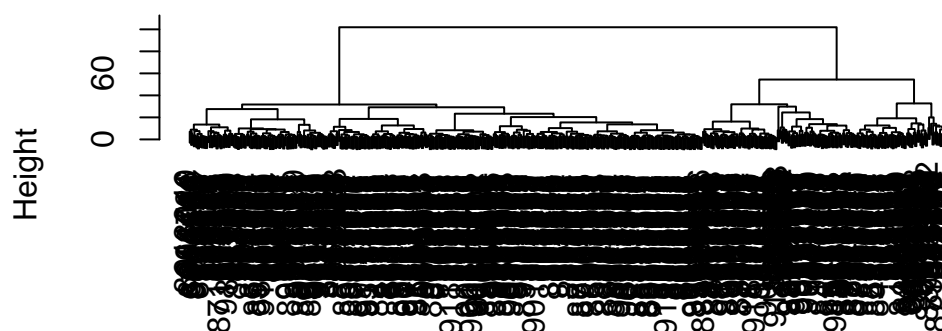
A: I prefer the complete method when trying hierarchical clustering as it compares the max values of each “cluster” in order to judge similarity, and I believe that’s a good way of comparison as it makes sure to contain all the data points into consideration.

Combining methods

Creating hierarchical clustering model using `method=“ward.D2”`.

```
wisc.pr.hclust <- hclust(data.dist, "ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



Checking the clustering:

```
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
grps  
  1   2  
184 385
```


Checking the two cluster groups in terms of diagnosis:

```
table(grps, diagnosis)
```

```
      diagnosis  
grps   B    M  
  1   20 164  
  2  337  48
```

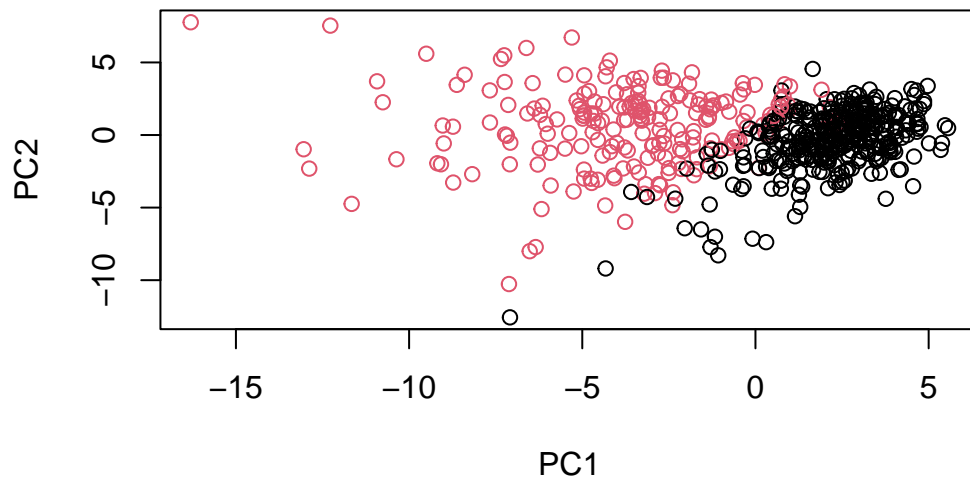
Plotting the principal components again categorized by the groups:

```
plot(wisc.pr$x[,1:2], col=grps)
```




Comparing it to the graph where it is categorized by diagnosis:

```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



Changing the graph that was categorized by the groups to follow the same color scheme as the PCA diagnosis graph by changing order of 1 and 2 in grps:

```
g <- as.factor(grps)
g <- relevel(g,2)
plot(wisc.pr$x[,1:2], col=g)
```

Using the distance along the first seven principal components for clustering so that we could compare:

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method="ward.D2")
```

cutting this hierarchical cluster into 2 clusters:

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

```
table(wisc.pr.hclust.clusters,diagnosis)
```

	diagnosis	
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

The newly created model helps separate the two diagnosis almost as well as the Euclidean distances taken from the data that was plotted with hierarchical clustering with the method = "ward.D2".

Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters,diagnosis)
```

```
          diagnosis
wisc.hclust.clusters  B  M
1      12 165
2       2   5
3    343  40
4       0   2
```


The hierarchical clustering models created in previous sections are not as effective as it causes clusters 2 and 4 to not contain enough data to consider the clusters to be significant.

Prediction

Predicting the diagnosis using the `predict()` function; new cancer data from two patients are downloaded and plotted against the PCA graph:

```
#downloading data
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)

#plotting with the PCA graph
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16: Which of these new patients should we prioritize for follow up based on your results?

A: We should prioritize patient 2 as it resides within the collection of data that were diagnosed as malicious.